

# Concurrency: Lock, Group and Data Races

23 March 2020

Ákos Frohner

# Overview

Part 1: Goroutines and Channels

Part 2: (this) Lock, WaitGroup and Data Races

Part 3: Context, Timeout, Rate Control

Examples are in the **ex?** subdirectories of

[github.com/golang-elte-2020-public/concurrency](https://github.com/golang-elte-2020-public/concurrency) (<https://github.com/hu-univ-golang/golang-elte-2020-public/tree/master/concurrency>)

## Last exercise: parallelize checksum calculation

```
// TODO: parallelize the checksum calculation
for _, path := range Files() {
    hash, err := Hash(path)
    if err != nil {
        fmt.Printf("ERROR: %s\n", err)
        continue
    }
    fmt.Printf("%x\t%s\n", hash, path)
}
```

## Solution:

```
type result struct {  
    path string  
    hash []byte  
    err  error  
}
```

```
files := Files()  
ch := make(chan *result, len(files))  
for _, path := range files {  
    go func(path string) {  
        hash, err := Hash(path)  
        ch <- &result{path: path, hash: hash, err: err}  
    }(path)  
}  
  
for range files {  
    r := <-ch  
    if r.err != nil {  
        fmt.Printf("ERROR in %q: %s\n", r.path, r.err)  
    } else {  
        fmt.Printf("%x\t%s\n", r.hash, r.path)  
    }  
}
```

## How to...

- store the result?
- handle unknown number of inputs?
- handle errors?

## Use case: Bazel

(re)building a project by tracking (file) dependencies

- run "go build ex3.go", when "ex3.go" changes
- change ~ different checksum

Continuously tracks and updates a path→hash map.

[bazel.build/](https://bazel.build/) (<https://bazel.build/>)

# Simplified example

Simplified environment to focus on the language constructs:

```
func Hash(path string) int {  
    time.Sleep(100 * time.Millisecond)  
    return len(path) // *not* collision free hash  
}
```

```
func Files() []string {  
    return []string{"ex1/ex1.go", "ex2/cksum.go", "ex3/ex3.go"}  
}
```

## Exercise 3: collecting results

```
// TODO: parallelize the calculation
results := make(map[string]int)
for _, path := range Files() {
    results[path] = Hash(path)
}
```

[play.golang.org/p/0yjPboEDb1k](https://play.golang.org/p/0yjPboEDb1k) (<https://play.golang.org/p/0yjPboEDb1k>)



# NOT a solution

```
// TODO: parallelize the calculation
ch := make(chan struct{})
results := make(map[string]int)
files := Files()
for _, path := range files {
    go func(path string) {
        results[path] = Hash(path)
        ch <- struct{}{}
    }(path)
}

for range files {
    <-ch
}
```

# Data race detection

```
$ go run -race no.go
=====
WARNING: DATA RACE
Write at 0x00c000078150 by goroutine 9: ...
    main.main.func1()
        .../concurrency/ex3/no.go:15 +0x69

Previous write at 0x00c000078150 by goroutine 7: ...
    main.main.func1()
        .../concurrency/ex3/no.go:15 +0x69

Goroutine 9 (running) created at:
    main.main()
        .../concurrency/ex3/no.go:14 +0x13d

Goroutine 7 (running) created at:
    main.main()
        .../concurrency/ex3/no.go:14 +0x13d
=====
map[ex1/ex1.go:10 ex2/cksum.go:12 ex3/ex3.go:10]
Found 1 data race(s)
```

# Mutex

Sometimes all you need is a lock:

- `sync.Mutex`: Lock, Unlock
- `sync.RWMutex`: Lock, Unlock, RLock, RUnlock

[golang.org/sync](https://golang.org/pkg/sync/) ([https://golang.org/sync](https://golang.org/pkg/sync/))

## Exercise 3...

## Exercise 3: solution with Mutex

```
// TODO: parallelize the calculation
ch := make(chan struct{})
mu := sync.Mutex{}
results := make(map[string]int)
files := Files()
for _, path := range files {
    go func(path string) {
        mu.Lock()
        results[path] = Hash(path)
        mu.Unlock()
        ch <- struct{}{}
    }(path)
}

for range files {
    <-ch
}
```

## Exercise 3: solution with channel

```
// TODO: parallelize the calculation
type result struct {
    path string
    hash int
}
ch := make(chan *result)
files := Files()
for _, path := range files {
    go func(path string) {
        hash := Hash(path)
        ch <- &result{path, hash}
    }(path)
}

results := make(map[string]int)
for range files {
    r := <-ch
    results[r.path] = r.hash
}
```

## How to...

- store the result ✓
- handle unknown number of inputs?
- handle errors?

## Exercise 4: wait for all results

The results are collected in a shared map.

```
// TODO: wait for all results
results := make(map[string]int)
mu := sync.Mutex{}
for _, path := range Files() {
    go func(path string) {
        mu.Lock()
        defer mu.Unlock()
        results[path] = Hash(path)
    }(path)
}
```

[play.golang.org/p/8Z6ELbwckkB](https://play.golang.org/p/8Z6ELbwckkB) (<https://play.golang.org/p/8Z6ELbwckkB>)



# Wait (group)

A thread-safe counter to wait until a group of goroutines finish.

```
var wg sync.WaitGroup
```

- `wg.Add(x)` to add `x` to the counter
- `wg.Done()` to decrement the counter
- `wg.Wait()` to block until the counter becomes zero

```
wg.Add(1)
go func() {
    defer wg.Done()
    ...
}()
wg.Wait()
```

## Exercise 4...

## Exercise 4: solution

```
// TODO: wait for all results
results := make(map[string]int)
mu := sync.Mutex{}
wg := sync.WaitGroup{}
for _, path := range Files() {
    wg.Add(1)
    go func(path string) {
        defer wg.Done()
        mu.Lock()
        defer mu.Unlock()
        results[path] = Hash(path)
    }(path)
}
wg.Wait()
```

## Range over channel revisited

Current code executes fix number of receives:

```
files := Files()
...
results := make(map[string]int)
for range files {
    r := <-ch
    results[r.path] = r.hash
}
```

We should read from the channel until it is closed:

```
ch := make(chan *result)
...
results := make(map[string]int)
for r := range ch {
    results[r.path] = r.hash
}
```

Exercise 5: close the channel!

## Exercise 5: wait for last send

```
// TODO: close the channel after the last send
type result struct {
    path string
    hash int
}
ch := make(chan *result)
for _, path := range Files() {
    go func(path string) {
        hash := Hash(path)
        ch <- &result{path, hash}
    }(path)
}
time.Sleep(200 * time.Millisecond) // to wait for all results
close(ch)                          // that "range ch" works

results := make(map[string]int)
for r := range ch {
    results[r.path] = r.hash
}
```

[play.golang.org/p/GcZ3pPiBd5b](https://play.golang.org/p/GcZ3pPiBd5b) (<https://play.golang.org/p/GcZ3pPiBd5b>)

## Exercise 5...

## Exercise 5: deadlock

```
ch := make(chan *result)
wg := sync.WaitGroup{}
for _, path := range Files() {
    wg.Add(1)
    go func(path string) {
        defer wg.Done()
        hash := Hash(path)
        ch <- &result{path, hash}
    }(path)
}
wg.Wait()
close(ch) // cannot close before last send

results := make(map[string]int)
for r := range ch {
    results[r.path] = r.hash
}
```

[play.golang.org/p/iTSdC0h\\_tKi](https://play.golang.org/p/iTSdC0h_tKi) ([https://play.golang.org/p/iTSdC0h\\_tKi](https://play.golang.org/p/iTSdC0h_tKi))

## Exercise 5: solution

```
ch := make(chan *result)
wg := sync.WaitGroup{}
for _, path := range Files() {
    wg.Add(1)
    go func(path string) {
        defer wg.Done()
        hash := Hash(path)
        ch <- &result{path, hash}
    }(path)
}
go func() {
    wg.Wait()
    close(ch)
}()

results := make(map[string]int)
for r := range ch {
    results[r.path] = r.hash
}
```



## How to...

- store the result ✓
- handle unknown number of inputs ✓
- handle errors?

## Hash may return an error

```
func Hash(path string) (int, error) {  
    time.Sleep(100 * time.Millisecond)  
    if len(path) == 12 {  
        return 0, fmt.Errorf("cannot calculate hash for %q", path)  
    }  
    return len(path), nil  
}
```

## Exercise 6: handle errors

Using a shared map for the results:

```
// TODO: do not print results in case of any error
results := make(map[string]int)
mu := sync.Mutex{}
wg := sync.WaitGroup{}
for _, path := range Files() {
    wg.Add(1)
    go func(path string) {
        defer wg.Done()
        if hash, err := Hash(path); err == nil {
            mu.Lock()
            defer mu.Unlock()
            results[path] = hash
        } else {
            fmt.Printf("ERROR %s\n", err)
        }
    }(path)
}
wg.Wait()
fmt.Println(results)
```

[play.golang.org/p/LtfgGhZOS1F](https://play.golang.org/p/LtfgGhZOS1F) (<https://play.golang.org/p/LtfgGhZOS1F>)

# WaitGroup with error handling

```
import "golang.org/x/sync/errgroup"
```

- `sync.WaitGroup()` → `errgroup.WithContext(ctx)`
- `wg.Add(1), go, wg.Done()` → `eg.Go(func() error)`
- `wg.Wait()` → `err := eg.Wait()`

```
eg, ectx := errgroup.WithContext(ctx)
eg.Go(func() error {
    ...
    return nil
})
err := eg.Wait()
```

[godoc.org/golang.org/x/sync/errgroup](https://godoc.org/golang.org/x/sync/errgroup) (<https://godoc.org/golang.org/x/sync/errgroup>)

## Exercise 6...

## Exercise 6: solution

```
// TODO: do not print results in case of any error
results := make(map[string]int)
mu := sync.Mutex{}
eg, _ := errgroup.WithContext(context.Background())
for _, path := range Files() {
    path := path
    eg.Go(func() error {
        hash, err := Hash(path)
        if err != nil {
            return err
        }
        mu.Lock()
        defer mu.Unlock()
        results[path] = hash
        return nil
    })
}
if err := eg.Wait(); err != nil {
    fmt.Printf("ERROR %s\n", err)
} else {
    fmt.Println(results)
}
```

## How to...

- store the result ✓
- handle unknown number of inputs ✓
- handle errors ✓

## Exercise 7: watching file changes

- periodically scan all files under a directory
- compare their checksum with the previous scan



## Exercise 7: Hash() returns Hashed

```
type Hashed struct {  
    Path string  
    Hash []byte  
    Err  error // Hash is invalid, in case of an error  
}
```

```
func Hash(path string) *Hashed {  
    f, err := os.Open(path)  
    if err != nil {  
        return &Hashed{Path: path, Err: err}  
    }  
    defer f.Close()  
  
    h := sha1.New()  
    if _, err := io.Copy(h, f); err != nil {  
        return &Hashed{Path: path, Err: err}  
    }  
    return &Hashed{Path: path, Hash: h.Sum(nil)}  
}
```

## Exercise 7: HashedEqual()

```
func HashedEqual(before, after *Hashed) bool {
    if before == nil || after == nil {
        return before == nil && after == nil
    }
    if before.Path != after.Path {
        return false
    }
    if be, ae := before.Err != nil, after.Err != nil; be || ae {
        return be == ae
    }
    if len(before.Hash) != len(after.Hash) {
        return false
    }
    for i := 0; i < len(before.Hash); i++ {
        if before.Hash[i] != after.Hash[i] {
            return false
        }
    }
    return true
}
```

## Exercise 7: CompareFilesets()

```
type FileSet map[string]*Hashed
```

```
func CompareFileSets(before, after FileSet) (added, edited, deleted []string) {  
    for bp, bh := range before {  
        switch ah, has := after[bp]; {  
            case !has:  
                deleted = append(deleted, bp)  
            case !HashedEqual(bh, ah):  
                edited = append(edited, bp)  
            }  
        }  
    for ap := range after {  
        if _, has := before[ap]; !has {  
            added = append(added, ap)  
        }  
    }  
    return added, edited, deleted  
}
```

## Exercise 7: periodically compare

```
func main() {  
    prev := HashAll()  
    for ts := range time.Tick(time.Second) {  
        curr := HashAll()  
        added, edited, deleted := CompareFileSets(prev, curr)  
        prev = curr  
        if len(added)+len(edited)+len(deleted) > 0 {  
            fmt.Printf("files have changed at %v\n", ts)  
            fmt.Printf("\tadded: %q\n", added)  
            fmt.Printf("\tedited: %q\n", edited)  
            fmt.Printf("\tdeleted: %q\n", deleted)  
        }  
    }  
}
```

## Exercise 7: checksum again

```
func HashAll() FileSet {  
    // TODO: parallelize the checksum calculation  
    results := make(FileSet)  
    for _, path := range Files() {  
        results[path] = Hash(path)  
    }  
    return results  
}
```

- store the results in a shared map
- wait until all results arrive

Thank you

Ákos Frohner

[szamcsi@gmail.com](mailto:szamcsi@gmail.com) (mailto:szamcsi@gmail.com)

